Hecho con Metabase

TEST DASHBOARD A

Overview

Case A evaluates a digital Tullave recharge workflow orchestrated in n8n. The experiment measures end-to-end latency, interoperability across integrated services, observed fault tolerance, and service-level timing under controlled execution scenarios. The results shown here must be interpreted as experimental evidence rather than as production performance.

How to read this dashboard

This dashboard must be read as a controlled experiment. The runner intentionally generated successful executions, business rejections, injected service failures, and latency-stressed runs. Therefore, non-successful executions do not represent a real production error rate; they represent the observed distribution under designed test conditions.

Total Executions

avg_latency_ms	p50	p95	p99
851.51	373	3,970.6	6,321.92

1 fila

Execution Status Distribution

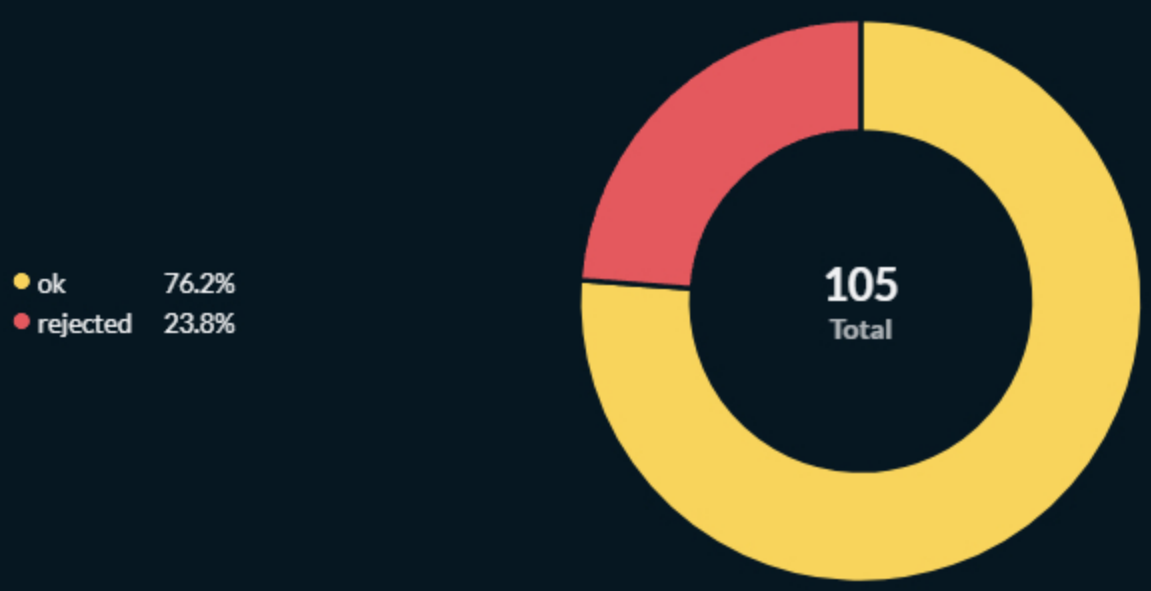

Experimental results

This dashboard must be read as a controlled experiment. The runner intentionally generated successful executions, business rejections, injected service failures, and latency-stressed runs. Therefore, non-successful executions do not represent a real production error rate; they represent the observed distribution under designed test conditions.

Failure Reason Compososition

Performance

End-to-end latency represents the total workflow processing time, while service latency reflects the observed contribution of each integrated component within the experiment. These values are useful for comparison inside the workflow, not as contractual SLA indicators.

End-to-End Latency Trend

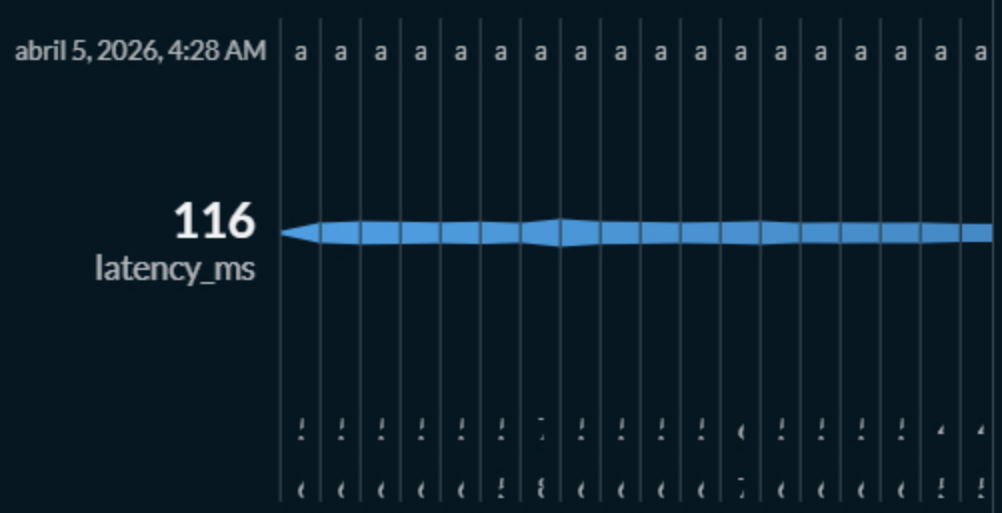

Average Latency by Service

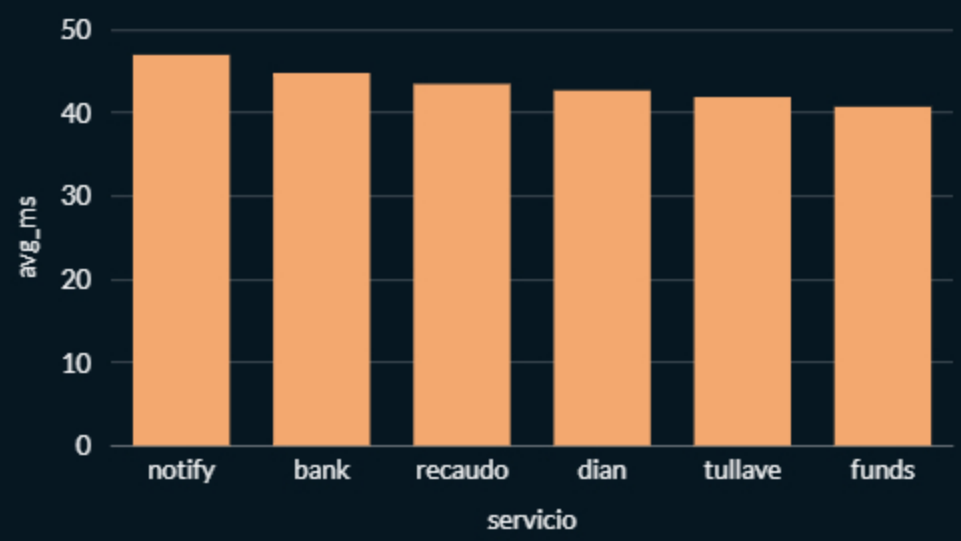

Interoperability

Interoperability in Case A must be interpreted as observed service participation and successful contractual interaction within the workflow. The dashboard therefore combines workflow-level completion metrics with service-level evidence such as latency presence, contractual conformance, and functional milestones.

Observed Workflow Interoperability

total_ejecuciones	ejecuciones_ok	ejecuciones_no_ok	tasa_exito_global_pct
105	80	25	76.19

1 fila

Observer Service Conformance

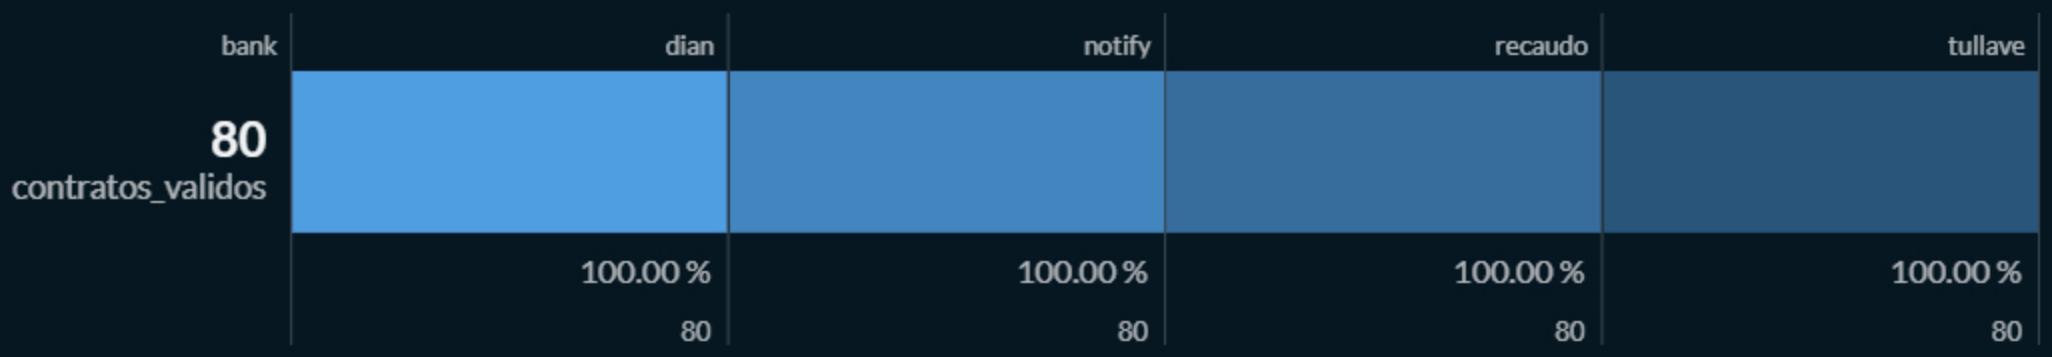

Complementary Technical Evidence

Resource Consumption Evidence

contenedor	cpu_avg_pct	ram_avg_mib	net_io_rx	net_io_tx	block_io_read	block_io_write
metabase	0.63	1,171.46	80 MB	22.8 MB	17.4 MB	901 MB
n8n	0.12	379.6	105 MB	44 MB	257 MB	50.5 MB
n8n-postgres	0.03	73.88	30.4 MB	178 MB	100 MB	29.2 MB
casea-mock-api	0.1	54.52	0.508 MB	0.539 MB	136 MB	0 MB
promedio	0.22	419.86	53.98 MB	61.33 MB	127.60 MB	245.18 MB

5 filas

Resource Consumption Evidence

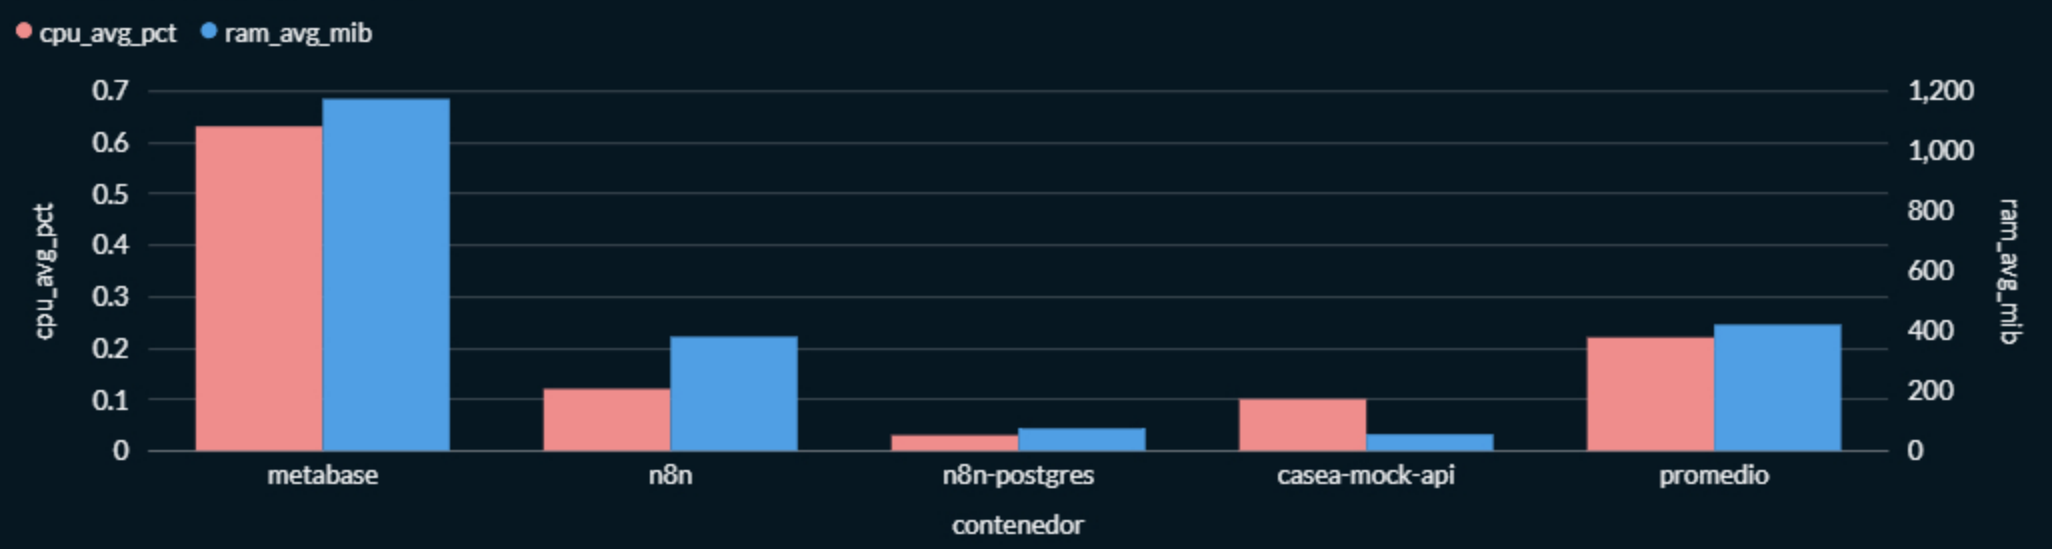

Transient Failure Recovery rate

This indicator shows the proportion of transient failures that were successfully absorbed through retry logic without ending in a final unrecovered error.

Transient Recovery Rate

total_transient	recovered_ok	not_recovered	recovery_rate_pct
14	14	0	100

1 fila

Fault Tolerance

The total dead-letter count shows how many executions ended in a technical fault, while the distribution by failed node shows where those faults concentrate inside the workflow.

Total Dead-Letter Events

total_transient	recovered_ok	not_recovered	recovery_rate_pct
14	14	0	100

1 fila

Dead-Letter Events by Failed Node

Workflow design

Workflow Desing Evidence

aspect	workflow_evidence
total nodes	25
external HTTP integrations	6
decision nodes	2
result persistence	3 inserts into PostgreSQL
terminal outputs	success, payment rejected, insufficient funds
traceability	transaction_id, x-correlation-id, status, fail_reason, start_time, end_time
service-level metrics	funds, bank, recaudo, tullave, dian, notify

7 filas

Enhanced Operational Detail

This table exposes execution-level evidence for the experiment. It provides traceability for individual transactions, including identifiers, functional status, fail reasons, timestamps, and service-level latencies

Operational Detail

transaction_type	status	latency_ms	service_latency_ms_funds	service_latency_ms_bank	service_latency_ms_recaudo
tx-transient-001	ok	1,745	172	44	25

tx-transient-001	ok	2,944	225	15	132
tx-transient-001	ok	4,101	171	17	186
tx-transient-001	ok	6,347	101	176	85
tx-transient-001	ok	4,102	99	162	50
tx-transient-001	ok	717	13	95	105
tx-transient-001	ok	2,959	52	133	245
tx-transient-001	ok	9,390	136	116	119

104 filas

Technical Evaluation Summary

Technical Evaluation

metric	indicator	result	status	observations
Data security and governance	Transaction-level traceability	Compliant	OK	End-to-end traceability is observed
Data security and governance	Credential exposure in logs	Not observed	OK	The table stores transactional data and metrics, not credentials
Fault tolerance	Events routed to dead-letter	Compliant	OK	Technical errors are captured and the failed node is identified
Fault tolerance	Automatic recovery from transient failure	Compliant	OK	Automatic recovery was evaluated using configured retries
Fault tolerance	Configured automatic retries	Compliant	OK	Retries were configured with up to 3 attempts per node
Resource consumption	CPU and RAM usage under nominal scenario	Compliant	OK	Observed resource consumption in the local environment was lo
Community and support availability	Documentation and ecosystem	Favorable	OK	This is a documentation-based metric, not derived from the trans
Workflow maintainability	Separation of responsibilities	Acceptable	OK	Separation improves maintainability and test execution
Local deployment feasibility	Reproducible self-hosted deployment	Compliant	OK	Local deployment was reproducible in the development environm

9 filas